

Towards P4 v1.2 - a proposal -

Mihai Budiu
Barefoot Networks
February 2016



Goals of 1.2

- Simplify and clean-up P4
- Provide a “stable” language version
 - I.e., we work very hard to keep the language **backwards compatible** from 1.2 onwards
- Provide a precise specification of the language
 - Including semantics of all constructs
- Provide a reference implementation of the language
 - Compiler front-end
 - Example programs
 - Behavioral simulator
- Address feedback received on P4 v1.0 and v1.1





P4 v1.2 vs 1.1



- An incremental evolution of P4 v1.1
- Same abstraction level
- Same core constructs
 - Parsers, control, match/action tables, actions, headers, metadata
- Same computational restrictions
 - No unbounded loops, no FP, no pointers, no recursion, constant work per header byte
- Simplified and clarified
- Avoid inventing new language constructs
 - Reuse well-understood tools and techniques as much as possible
- Prepare language for future evolution through *growth*
 - Architectural features caused most of P4's growth

Language clean-up

- Break language into three parts
 - Core language (part of the language spec)
 - Packet processing language
 - Language constructs to describe architectures
 - Standard library (e.g., common to *all* architectures)
 - A standard architecture spec
 - Prototypes for architectural blocks and intrinsic metadata
 - Library with extern blocks declarations (e.g., checksums)
- Write a specification for the control/data-plane API
- Libraries and architectures are written in P4
- These separate specifications evolve independently
- Architecture evolution becomes much easier



Desirable P4 v1.2 features

- Strong static typing
- Simpler syntax
- Few undefined behaviors
- No runtime exceptions/traps
- Explicit deparser specification
- Clear evaluation results (“declarative” => “deterministic”)
- Lexical scoping
- Support for writing modular programs
- Support for error handling
- Parameterization (e.g., “how many bits to specify an output port?”)
- Compile-time resource allocation (e.g., checksum units, tables, etc.)
- Simple extensibility hooks (e.g., Java-like annotations)



Details for some
proposed constructs



Architecture specification




A proposal for architecture specification was given in December 2015

- That presentation is included as an appendix
- The architectural specification language included the following features:
 - struct/header types
 - parsers/control/packages architectural blocks
 - prototypes for architectural blocks
 - generic types (templates)
 - parameterized architectural blocks
 - separation of declaration vs. instantiation

Interaction with architecture

- Intrinsic metadata
 - Action occurs at the “end” of the pipeline
- Extern object method invocations
 - Action occurs instantly
- No “delayed” execution
 - E.g., drop, field_list_calculation, generate_digest
 - Order of delayed executions was unspecified
 - Order of side-effects and delayed executions was unspecified
- The meaning of a P4 program should be unambiguous



Moving constructs from P4 to libraries

- Custom primitive actions declarations
- `field_list_calculation` (e.g., checksums, `modify_field_with_hash_based_offset`)
- `parser_value_set`
- `generate_digest`
- cloning, recirculation, resubmission, mirroring
- Counters, meters, registers
- Action profiles
- Saturated types
- In general, all constructs which “look” non-portable across all architectures




Explicit packet (in standard library)

```
extern packet_in {  
    void extract<T>(out T hdr);  
    void extract<T>(out T varSizeHeader, in bit<32> size);  
    T lookahead<T>();  
}  
  
extern packet_out {  
    void emit<T>(in T hdr);  
}  
  
parser prs(packet_in p, Headers h) {  
    p.extract(h.eth);  
}
```



Deparsers

- In P4 v1.1
 - Sometimes impossible to infer
 - Users have no control
 - Hacks for creating fabric headers
- In P4 v1.2
 - Just another control block
 - Should clearly specify sequence of actions (emit, checksums)




```
control deparser(in headers h, packet_out p)
{
    Checksum16() ck;
    apply {
        ck.clear();
        h.ip.hdrChecksum = 0;
        ck.update(h.ip);
        p.emit(h.ethernet);
        h.ip.hdrChecksum = ck.get();
        p.emit(h.ip);
    }
}
```

Parameterization support

- Writing portable and modular programs
- **typedef**
 - E.g., **typedef** bit<8> Port_t;
- **enum**
 - E.g., **enum** ChecksumType { crc16, crc32 }
- constant declarations
 - **const** Port_t CPU_PORT = 16;
- Generics (templates)
 - E.g., parser<H>(packet_in p, out H headers)
- Constructors
- (See also the architectural description proposal)





Simpler syntax

- `modify_field`, `set_metadata` => assignment statements
 - `modify_field(a, b)` => `a = b`
 - `set_metadata(a, b)` => `a = b`
- Add a few useful operators: masking, concatenation, bit selection, mux, range
- Convert keywords to methods or fields
 - `valid(a)` => `a.valid`
 - `add_header(a)` => `a.setValid(true);`
 - `remove_header(a)` => `a.setValid(false);`
 - `copy_header(a, b)` => `a = b`
 - `hs[last]` => `hs.last`
 - `push(hs, 2)` => `hs.push_front(2)`


Richer type system

- **enum**
- **error**
- **header/struct**
- **header_union**
- Stacks of unions => option parsing
- Typed architecture blocks
 - parsers, control blocks, packages
 - (See also the architectural description proposal)



Extensibility hooks

- `@annotation(expression)`
- Allows for some language evolution without spec changes
- Pragma-like
- Typed
- Apply to specific language elements
- Similar to Java `@` annotations and C# [Attributes]
- Some annotations could become part of standard



Scoping

- Create lexical scopes
- Remove global variables
- Introduce local variables and parameters
 - Clarifies scope of intrinsic metadata
(See also the architectural description proposal)
 - Enables modular programs
- Declarations must precede uses (except parser states)



Error handling



- Accommodate various architectural constraints
 - (E.g., encoding of error codes)
- Add an **error** type (special enum-like type):
error { IncorrectVersion, HeaderTooShort }
- Parser exceptions => “reject” parser state
- Introduce an “assert” method, usable in parsers
`assert(h.ip.version == 4, IncorrectVersion);`
 - Assertion failure sets error code and transitions to the reject state
- Expose errors explicitly to control blocks
control ingress(**in error** parser_error, ...)

Flexible control-flow

- Control blocks:
 - Add a **return** statement
 - Add an **exit** statement
- Parsers
 - rename “**return**” to “**transition**”



To be continued...

- We will produce the following:
 - Draft design with all of these features
 - A “migration guide” mapping P4 v1.1 constructs to P4 v1.2
 - Example programs and program fragments
 - A written P4 v1.2 specification draft




Appendix

- The following slides include for reference the presentation from December 2015 on a proposal for describing architectures in P4

Abstracting switch architectures - a proposal -

November 30, 2015

The P4 tension



Universal

Customizable

P4 v1: Fixed Abstract Forwarding Model

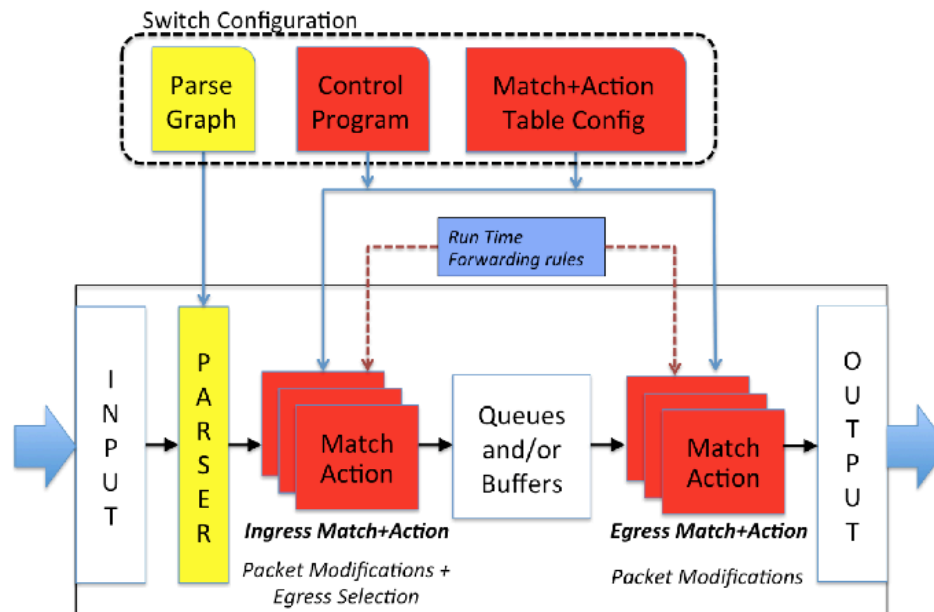
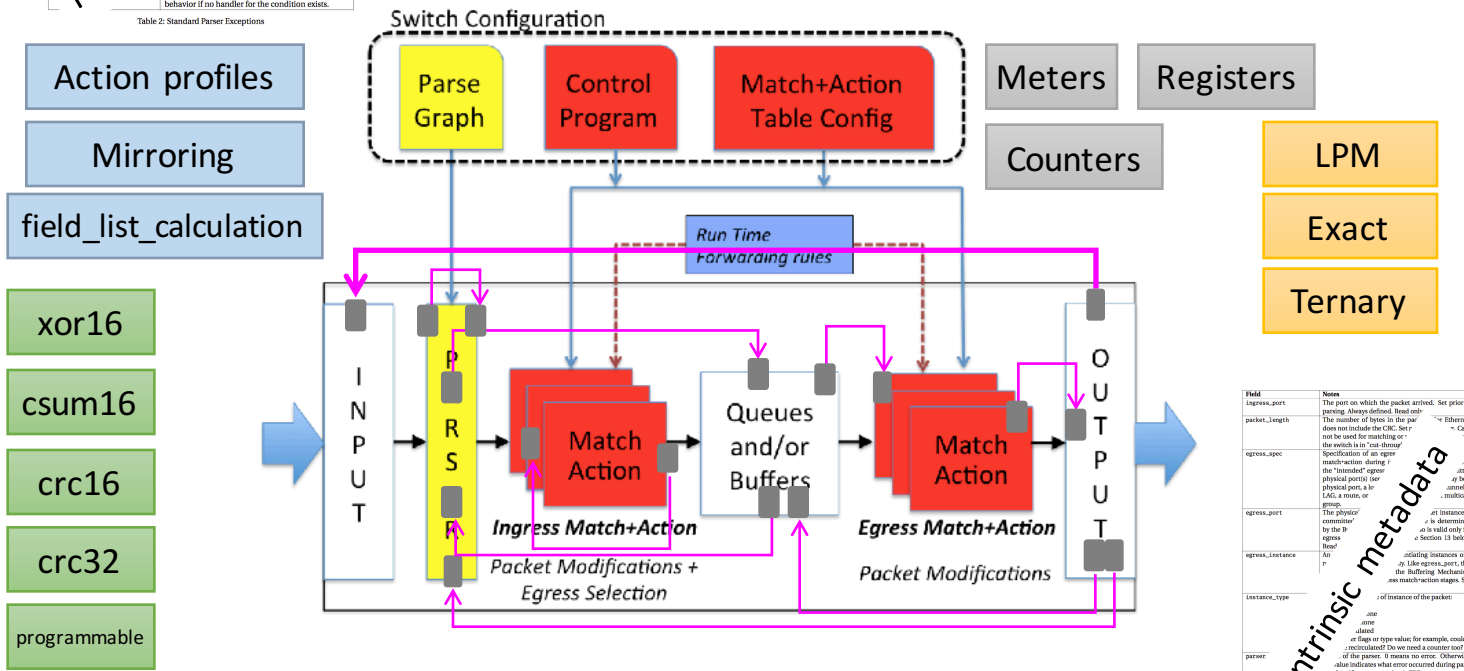


Figure 1: Abstract Forwarding Model

P4 v1: Details

Identifier	Exception Event
p4_pe_index_out_of_bounds	A header stack array is declared bound.
p4_pe_out_of_packet	There were not enough bytes to complete an action.
p4_pe_header_too_long	A header field was declared to be longer than the declared length.
p4_pe_header_too_short	A header field was less than the minimum length portion of the header.
p4_pe_unsupported	An action had no default specified but the default value was not in the case list.
p4_pe_parse_error	A parse error was detected.
p4_pe_match_error	This is not an exception itself, but allows the programmer to define a handler to specify the default behavior if no handler for the condition exists.

Table 2: Standard Parser Exceptions

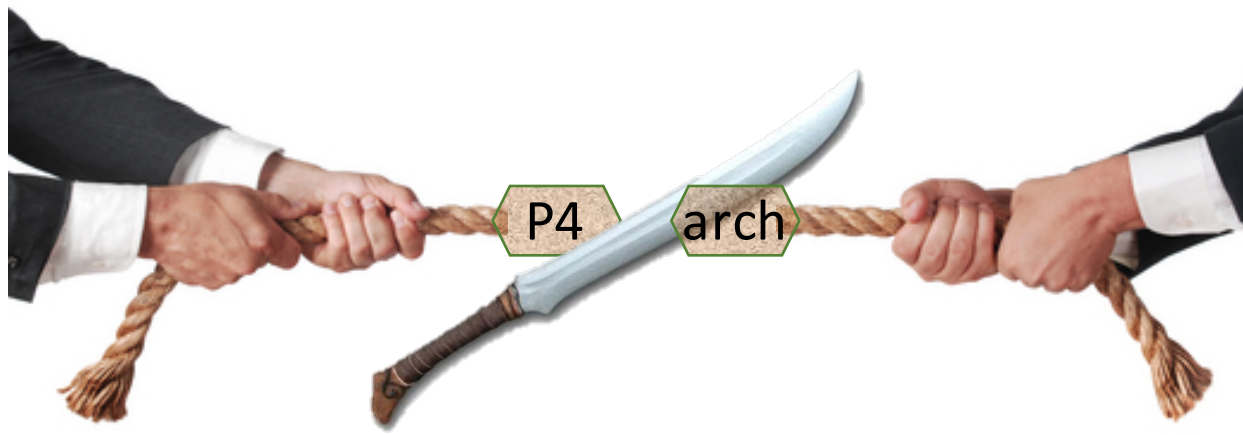


Field	Notes
ingress_port	The port on which the packet arrived. Set prior to parsing. Always defined. Read only.
packet_length	The number of bytes in the packet. Does not include the CRC. Set prior to parsing. Read only. Cannot be used for matching.
egress_spec	Specification of an egress match-action during parsing. It is the "intended" egress physical port, a logical port, a MAC, a VLAN, a tunnel, or a group.
egress_port	The physical port that the packet is committed to. It is determined by the ingress port and the egress_spec.
egress_instance	An instance of a match-action. It is valid only for the egress port.
instance_type	The type of the instance of the packet.
parser	The parser that parsed the packet. Do not use a counter for this field. It is not an intrinsic metadata field. It is a parser error occurred. This is an indication of the location in the parser program where the error occurred. Specific representation is TBD.

Intrinsic metadata

Table 3: Standard Intrinsic Metadata Fields

Divide and conquer

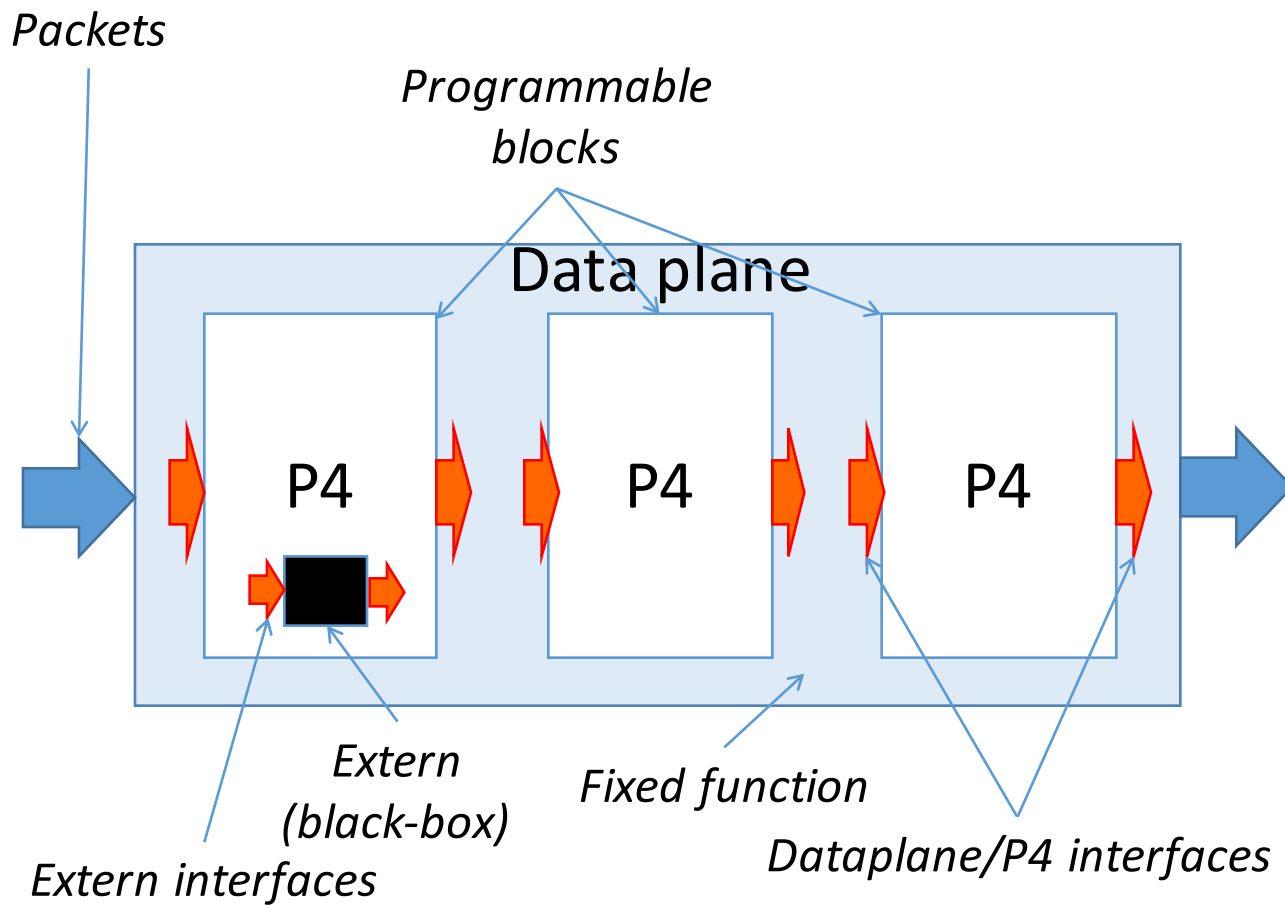


- Separate language definition from architecture definition
- Evolve them independently




Specifying architectures

- A device model describes what parts of a forwarding device can be programmed in P4.
- Each manufacturer can publish custom device models.
- The community defines a standard switch model for portability.
- Even if without custom switch models, this approach is useful, because it decouples the language evolution from the model evolution: new versions of the standard switch model do not require changing the language.

Generic Programmable Dataplane Model



P4 Support for multiple architectures



Inventing new language constructs

- Don't
- You will get them wrong
- Reuse constructs from other languages

- How would I do this in Java/C++?

Switch architecture in C++

```
// switch.hpp: written by manufacturer
```

```
struct MetaIn { int inputPort; }
```

```
struct MetaOut {  
    int outputPort;  
    bool drop;  
}
```

```
template<class T> class switch
```

```
{
```

```
    virtual void parser(const packet &p, T& headers)=0;
```



```
    virtual void control(T& headers,  
                        const MetaIn &in, MetaOut& out)=0;
```

```
}
```

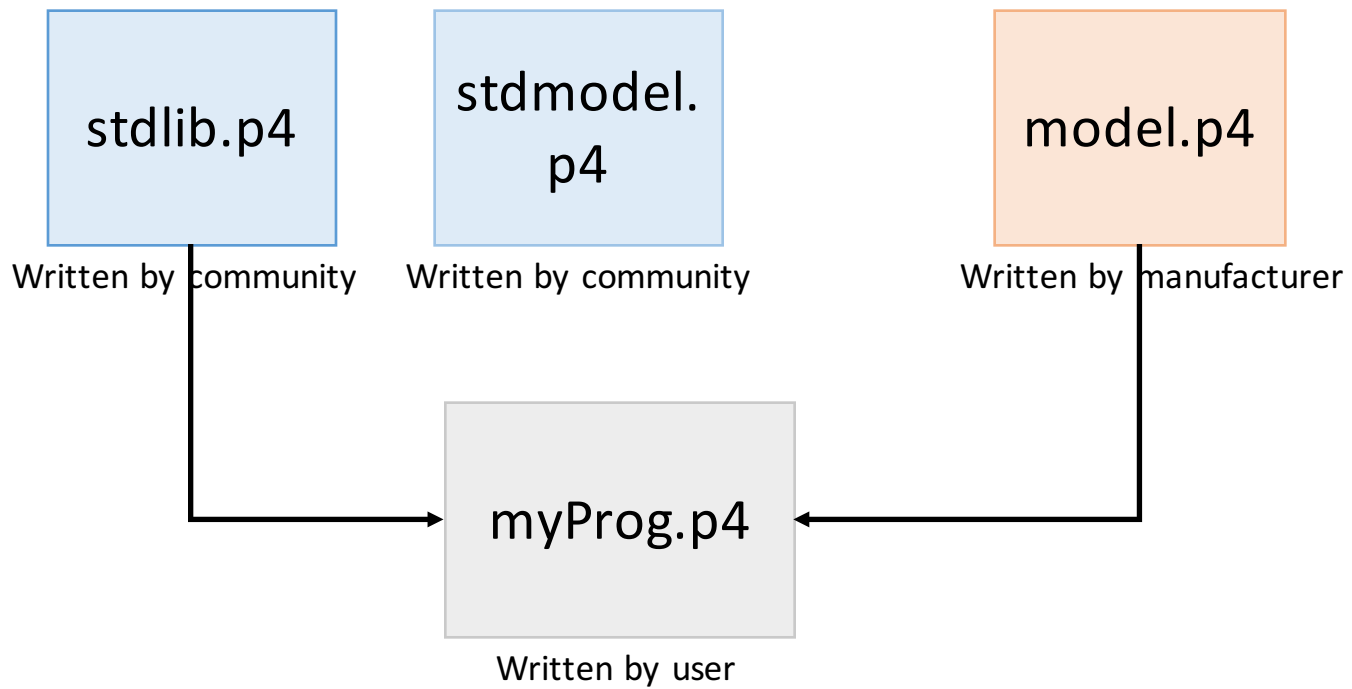
abstract methods = implemented by user

user-defined metadata

target-defined metadata



Structure of a P4 program



Detailed Design

How is this different from “whitebox”?

- This is a revision of the previous whitebox proposal
- Accomplishes same goals
- Slightly different approach
 - Break out whitebox into multiple simple constructs
 - parser, control, package
 - Allows for separate type-checking
 - Modeled after C++/Java OO
 - We can provide an operational semantics for all these constructs

P4 program skeleton

```
// standard definitions
#include "stdlib.p4"
// architecture description; includes Switch decl.
#include "arch.p4"
// user code
parser myParser ...
// architecture instantiation
Switch(myParser(), myControl()) main;
```

Preliminaries

- Add a proper “struct” type
- Can be used for metadata (replacing the **metadata** keyword)
- The **header** type is just for headers
- Structs can be nested (but not headers)

```
header ethernet { ... }
```

```
header ipv4 { ... }
```

```
struct headers_t { ethernet e; ipv4 ip; ... }
```

Basic building blocks

- **parser** and **control**
- They look like functions
 - Local scope
 - Arguments with directionality
 - They are typed
- **Rationale:**
 - **in** and **out** arguments indicate the scope of metadata and the user data. For example, the parser metadata cannot be accessed in the Ingress block.
 - Signatures allow type checking
 - Help with resource allocation, by delimiting the scope of various structures.

Parsers

- Rename **parser** -> **state**
- Use **parser** for grouping states

```
parser P(in parser_metadata_t pm,  
         out header_t headers)  
{  
    state start { ... extract(headers.e); ... }  
    state parse_ip { ... }  
}
```

Control blocks

```
control Ingress(inout headers_t headers,  
               in control_metadata_in cmi,  
               out control_metadata_out cmo)  
{  
  table t { ... }  
  action a { ... }  
  apply { /* control body */ }  
}
```

Declarations

- Architecture declares **prototypes** for programmable blocks
- Users define blocks with matching prototypes

```
parser P<H>(in parser_metadata pm, out H headers);  
control Ingress<H>(inout H headers,  
                    in control_metadata_in cmi,  
                    out control_metadata_out cmo);
```

- **Rationale:**
 - type variables indicate user-specified types
 - Type variables are only allowed in architectural specifications
 - users cannot write code containing type variables

Persistent Resources

- Compiler must allocate resources
- E.g., extern objects, tables, and blocks containing such objects
- Parsers and control blocks are persistent resources

```
parser name(arguments)  
{  
  stateful_Instantiations  
  state { ... }  
}
```

```
control name(arguments)  
{  
  stateful_Instantiations  
  apply { /* control flow here */ }  
}
```


Instantiating a resource

```
extern Checksum16 { ... }
```

```
parser MyParser(...)
```

```
{
```

```
    Checksum16 ck; // checksum unit instantiation
```

```
    ...
```

```
    state start { ... }
```

```
    state ipv4 { ... ck.verify(h.ipv4); ... }
```

```
}
```

Types and instances

- **parser** and **control** block declare types
- Types must be instantiated to be used

```
control IPv4Control(inout Headers headers)
{ ... }
```

```
control Ingress(inout Headers headers, ...)
{
    IPv4Control() ipv4control; // instantiate control block
    table acl { ... } // table instantiation
```

```
    apply {
        ...
        ipv4control.apply(headers); // invoke control instance
        acl.apply(); // invoke table instance
    }
}
```

Rationale for instantiations

- Parsers and control blocks are similar to classes in OO languages.
- Separating type declaration from instantiation allows one type to be instantiated multiple times.
- E.g,: configure a switch that has 4 ingress pipelines where each of them can be programmed independently: the programmer can write one type, and instantiate it 4 times, once for each pipeline.
- Instantiation is denoted using constructor invocation.

Packages

- A **package** is a container which may contain other packages, parsers and control blocks.
- The toplevel forwarding element is declared as a package by the architecture manufacturer and instantiated by the user.

Switch instantiation outline

```
// Architecture declaration by manufacturer
parser Parser<H>(out H headers, ...);
control Ingress<H>(inout H headers, ...);
control Deparser<H>(in H headers, ...);
// toplevel element:
package Switch<H>(Parser<H> p,
                  Ingress<H> ingress,
                  Deparser<H> deparser);
```

```
// Program written by user
struct head { ... }
parser MyParser(out head h, ...) { ... }
control MyIngress(inout head h) { ... }
control MyDeparser(in head h...) { ... }
```

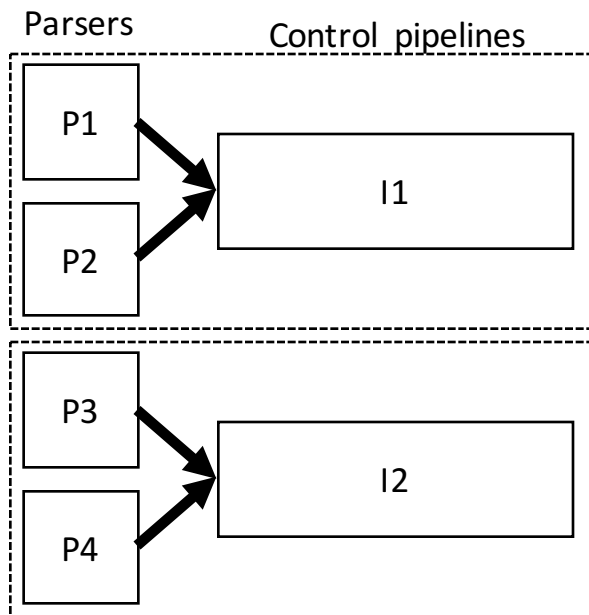
```
// toplevel element instantiation
Switch(MyParser(),
      MyIngress(),
      MyDeparser()) main;
```

H = **struct** head –
inferred by compiler

Rationale

- The manufacturer can specify complex switches, with many programmable surfaces.
- The type parameters allow various switch components to be linked with each other
 - (e.g., the headers from the parser are the input/output of the ingress pipeline and the input to the deparser
 - the user cannot write an ingress pipeline that accidentally processes different headers from the parser).
- The manufacturer can expose multiple switch models, and the user can choose which one to instantiate (e.g., a standard model, or a model with additional features).
- The user can explicitly instantiate each programmable surface of the switch with the desired implementation.

A Complex Example



```
parser P<H>(out H headers);  
control I<H>(inout H headers, ...);
```

```
package module<H>(P<H> p1, P<H> p2, I<H> pipe);
```

```
package switch<H1, H2>(module<H1> mod1,  
                      module<H2> mod2);
```

Different headers allowed

Parameterization

- Third parties can write pre-packaged P4 code, which can be reused in a modular way.
- To suit the needs of arbitrary users, these blocks may be parameterized
- Similar to C++/Java/ML Functors

```
control ACLControl(bool largeSize) // parameters  
                (inout Headers h) // arguments
```

```
{ ... }
```

```
#include "aclControl.p4"
```

```
control Ingress(inout Headers headers)
```

```
{
```

```
    ACLControl(false) aclControl; // instantiate with largeSize=false
```

```
    apply {
```

```
        aclControl.apply(headers);
```

```
    }
```

```
}
```